MOVEINSYNC ASSIGNMENT / FORWARD DEPLOYED ENGINEER

Where the 1-star reviews come from, and what to fix first

Prepared by Shreya S (reg. no. 12320371), 12 October 2026. Public sources only; no MoveInSync internal data was used.

Recommendation in brief

The employee app's rating is not low across the board. About 114K ratings average roughly 4.31, and 12.3% of them are 1-star while only 2.7% are 2 or 3-star. The score is pulled down by a specific block of 1-star reviews, so the job is to find and fix that block rather than polish the whole app.

In the small sample I could read, critical reviews cluster at **gate moments**, where the app stands between the employee and the cab: login OTP, the boarding OTP behind a seat-belt step, booking status that contradicts itself, and a no-show that arrives as a surprise. I recommend fixing these in order, starting with the boarding OTP because it touches safety, and running a weekly loop that turns each complaint cluster into an owned, tested fix.

The evidence is thin (13 Android reviews, 8 negative). Treat the pattern as a lead to confirm with internal data in the first three days, not as a proven result.

1. What the public data says

Fact. Official listing: MoveInSync, package com.moveinsync.ets, publisher MoveInSync Technology Solutions, updated 8 Oct 2026. The driver, field and workplace apps are separate listings. The Play page showed 4.6; AppBrain reports 4.32 from 114,607 ratings. Confirm in Play Console which figure leadership tracks.

Stars	Ratings	Share
5	82K	71.8%
4	15K	13.1%
3	2.2K	1.9%
2	0.96K	0.8%
1	14K	12.3%

Source: appbrain.com histogram (third party, rounded); shares computed by me.

App	AppBrain avg	1-star share	Ratings
MoveInSync	4.32	12.3%	114.6K
Safetrax Commuter	4.21	12.1%	24.1K
Routematic	3.81	17.9%	34.8K
eFmFm	3.50	35.1%	4.9K

MoveInSync is already close to the best of the comparable apps. Illustrative arithmetic, not a forecast: if half of today's 1-star ratings became 4-star, the average would move from about 4.31 to about 4.50.

2. What employees are complaining about (finding, small sample)

Pattern	Evidence (paraphrased, 2026)	Stage	Count	Conf.
Boarding OTP gated and offline-fragile	1-star, 30 Sep: seat-belt check hides the OTP; app needs network on each open. Help Center confirms the OTP shows after the seat-belt step.	Boarding	1 of 8	Low-Med
Booking state unclear	Play, 7 Jul: shuttle booking became two steps, reviewer thought a cab was booked. 1-star, 6 Oct: booked trip, push says not booked.	Booking	2 of 8	Medium
No-show arrives as a surprise	Play, 29 Jun (6 helpful votes): missed check-in after a late cab is marked no-show.	Pickup	1 of 8	Low-Med
Login OTP not received	1-star, 24 Sep: repeated tries, no OTP. Help Center and a public forum post describe the same.	Onboarding	1 of 8	Medium
Tracking and ETA inaccurate	Play, 1 Jul: slow GPS, wrong pickup and drop times. Peer app reviews report the same.	Live trip	1 of 8	Medium
Outside the app	1-star, 14 and 18 Sep: pickup location, routing and vendor service.	Operations	2 of 8	Medium

A crash complaint also appeared (1-star, 18 Sep, plus iOS reports after v10.42.x). Counts overlap. Sample: 3 Play listing reviews and 10 AppBrain comments, English only. No developer replies or app versions were visible.

Hypothesis (untested). A gate failure is silent and the employee often has no other route or choice of app, so the Play review becomes their only complaint channel. It would be disproved by internal data showing gate failures are rare per trip, or by 1-star reviewers mostly citing routing and vendor service rather than app steps.

3. What to fix, in order

Score = Frequency x Severity x Confidence / Effort, each rated 1 to 5. The ratings are my provisional judgement on a small sample. Safety-linked items go first whatever the score.

#	Action	F	S	C	E	Score	When
1	Boarding gate: OTP always reachable, trip cached offline, seat-belt prompt kept but never blocking	3	5	3	3	15	Now, safety first
2	Booking state: one clear confirmed status, remove the contradictory 'not booked' push	3	4	3	2	18	Now
3	Login OTP fallback: resend timer, alternate channel, helpdesk link	2	5	3	2	15	Now
4	No-show warning when the cab reaches the pickup point; show the check-in window	2	4	3	2	12	Medium term
5	Stale-ETA indicator ('updated 40s ago') instead of a confident wrong time	3	3	3	3	9	Medium term
6	Crash and ANR instrumentation; release notes that name what was fixed	2	4	2	3	5	Start week 1

Row 2 scores higher than row 1, but boarding goes first because the OTP and seat-belt step is a safety control and a blocked boarding is the worst moment for a trip to fail. I do not suggest removing the seat-belt check: the Help Center says the driver also records seat-belt status, so the fix is to keep data capture while making sure one failing step cannot stop boarding. This needs sign-off from whoever owns safety compliance.

Longer term: fix the GPS and ETA pipeline once internal telemetry shows where it drifts, and send routing and vendor complaints to the transport desks with a visible owner.

4. Ticket an engineer can start tomorrow: BOARD-1

Problem	A review says the seat-belt step hides the OTP and the app needs a network call every time it opens, which fails underground.
Evidence	Android 1-star dated 30 Sep 2026. Help Center: the OTP shows only after the seat-belt action, and the popup repeats at intervals.
Reproduce (to confirm)	Test tenant with seatBeltNotificationEnabled on. Open an assigned trip, tap View OTP, then repeat in airplane mode and on a throttled network.
Expected	The last synced trip and OTP show with a 'last updated' time. The seat-belt prompt takes one tap, and if it cannot load, the OTP still shows.
Acceptance	OTP visible within 2 taps. Offline open shows the cached trip. The seat-belt answer is still logged. The safety owner signs off the fallback wording.
Instrumentation	Events: otp_view_tapped, seatbelt_prompt_shown, seatbelt_prompt_answered, otp_displayed, time_to_otp_ms, offline_open, cache_hit.
Test cases	1) Prompt service times out, OTP still shows. 2) Airplane mode, cached trip opens. 3) Prompt answered, event logged once. 4) Cache older than the trip, stale banner shows. 5) Trip changed on the server, app refreshes on reconnect.

Written from public evidence. The flag name comes from the public Help Center; event names are proposals to match to the real codebase.

5. 30-day plan

When	What happens	Owner (role)
Days 1 to 3	Request Play Console reviews and ratings, crash data, OTP and booking event logs. Confirm which rating figure counts.	FDE with product analyst
Days 4 to 7	Check each hypothesis against internal data. Start crash and event instrumentation. Agree the seat-belt fallback with safety.	FDE, mobile lead, safety owner
Week 2	Build BOARD-1, the booking status fix and the login fallback. Test on a staging tenant.	Mobile and backend engineers, QA
Week 3	Staged rollout to a share of employees at willing customers, with a holdout. Reply to affected reviews.	Mobile lead, support lead
Week 4	Read the results, keep or roll back, write the next cycle from the first weekly report.	FDE, product manager

Roles are generic because I do not know MoveInSync's team structure. If internal data shows gate failures are rare, the order of the fixes changes before engineering time is spent.

6. How success is measured

Public metrics (available now)	Internal metrics (data requests)
Average rating and weekly review volume. Share of 1-star reviews. Reviews tagged to each journey stage. Complaint recurrence by app version once versions are visible.	Time to OTP displayed, OTP and booking failure rates, crash-free sessions, support contacts per trip, issue resolution time.

First experiment (BOARD-1). Primary outcome: p90 time to OTP displayed and share of boardings with an OTP failure. Guardrails: seat-belt confirmation rate, crash-free sessions, boarding support contacts. Comparison: staged rollout against a holdout of employees in the same company for 2 weeks. Roll back if seat-belt confirmation falls by more than 2 points (a provisional threshold to be set with the safety owner on the real baseline). I do not promise a star increase: ratings lag, reviews over-represent angry users, and review volume changes for unrelated reasons, so the star average is a lagging check, not the test.

7. A weekly loop, with a first prototype

Collect reviews from a Play Console export or API (a data request) plus a dated listing snapshot; deduplicate and detect language; tag stage and failure with rules, using a model only for unclear or mixed-language reviews, with a human check; build a fingerprint per cluster (stage, failure, version, recurrence, severity); draft a QA test and ticket with an owner; compare before and after the next week. No undocumented review endpoints are used.

Prototype. review_triage.py (rules only, no model calls) reads a review CSV, tags stage and failure, skips 4 and 5-star praise, and writes a tagged CSV, a report and draft QA test stubs. On the 13-review sample it finds 8 negatives: booking state unclear 2, routing or vendor 2, and one each for no-show, ETA, boarding OTP, login OTP and crash. **It is unvalidated:** the rules were written while reading these same reviews. A first run wrongly counted a praise review as a complaint, which I fixed.

How I got here: method and tools

Starting point. The brief gave no data, so I started from the official Google Play listing for the employee app and checked that it was the employee app and not the driver, field or workplace apps.

- Read the listing and the public review pages I could reach: the Play listing (3 visible reviews), AppBrain (rating histogram and 10 recent comments) and MoveInSync's own Help Center.
- Computed the average and star shares from the histogram with a short Python script, and did the same for three comparable commute apps (Safetrax, Routematic, eFmFm).
- Tagged the 13 Android reviews I could read (8 negative) by journey stage and failure type. Paraphrases are in reviews_sample.csv.
- Turned the pattern into a ranked fix list, a first ticket, a 30-day plan, a measurement and experiment plan, and a weekly process.
- Wrote and ran a small rules-only Python script that reproduces the tagging and drafts QA test stubs.

Tools. Web search and page fetches, Python, and ChatGPT and Claude (AI assistants) for reading, structuring and drafting. I reviewed the output and the conclusions are mine to defend.

Limits. The sample is small and not random, so frequencies are directional. I saw no developer replies or app versions. The listing showed 4.6 while the histogram implies about 4.31. The scoring is my judgement. The script was tuned on the same 13 reviews. Assumptions: employees often cannot choose another app, and a review mention is not the same as how often a failure happens.

First thing I would do with access. Pull the Play Console review export and the OTP and booking logs, then test whether the boarding, booking and login gates really account for the 1-star block.

Sources: play.google.com/store/apps/details?id=com.moveinsync.ets; appbrain.com pages for MoveInSync, Safetrax Commuter, Routematic and eFmFm (10 Oct 2026); helpcenter.moveinsync.com (seat-belt compliance and mobile app issues articles).